

```
In [53]: # Predicting players rating
# In this project you are going to predict the overall rating of soccer player
# based on their attributes such as
# 'crossing', 'finishing etc.
# The dataset you are going to use is from European Soccer Database(https://www.kaggle.com/hugomathien/soccer) has more
# than 25,000 matches and more than 10,000 players for European professional s
#occer seasons from 2008 to 2016.

# About the Dataset
# The ultimate Soccer database for data analysis and machine learning The data
# set comes in the form of an SQL database
# and contains statistics of about 25,000 football matches, from the top footb
# all league of 11 European Countries.
# It covers seasons from 2008 to 2016 and contains match statistics (i.e: scor
# es, corners, fouls etc...) as well as the
# team formations, with player names and a pair of coordinates to indicate the
# ir position on the pitch. +25,000
# matches +10,000 players 11 European Countries with their Lead championship S
#easons 2008 to 2016 Players and
# Teams' attributes* sourced from EA Sports' FIFA video game series, including
# the weekly updates Team line up with
# squad formation (X, Y coordinates) Betting odds from up to 10 providers Deta
# iled match events (goal types, possession,
# corner, cross, fouls, cards etc...) for +10,000 matches The dataset also has
# a set of about 35 statistics for each
# player, derived from EA Sports' FIFA video games. It is not just the stats t
# hat come with a new version of the game
# but also the weekly updates. So for instance if a player has performed poorl
# y over a period of time and his stats get
# impacted in FIFA, you would normally see the same in the dataset.
```

```
In [1]: import sqlite3
import pandas as pd
import numpy as np
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from math import sqrt
import warnings
warnings.filterwarnings('ignore')
```


```
In [2]: # Create your connection.
cnx = sqlite3.connect('database.sqlite')
df = pd.read_sql_query("SELECT * FROM Player_Attributes", cnx)
```

```
In [3]: df.head()
```

```
Out[3]:
```

	id	player_fifa_api_id	player_api_id	date	overall_rating	potential	preferred_foot	attackin
0	1	218353	505942	2016-02-18 00:00:00	67.0	71.0	right	
1	2	218353	505942	2015-11-19 00:00:00	67.0	71.0	right	
2	3	218353	505942	2015-09-21 00:00:00	62.0	66.0	right	
3	4	218353	505942	2015-03-20 00:00:00	61.0	65.0	right	
4	5	218353	505942	2007-02-22 00:00:00	61.0	65.0	right	

5 rows × 42 columns



```
In [5]: df.shape
```

```
Out[5]: (183978, 42)
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 183978 entries, 0 to 183977
Data columns (total 42 columns):
id                183978 non-null int64
player_fifa_api_id 183978 non-null int64
player_api_id     183978 non-null int64
date              183978 non-null object
overall_rating    183142 non-null float64
potential         183142 non-null float64
preferred_foot    183142 non-null object
attacking_work_rate 180748 non-null object
defensive_work_rate 183142 non-null object
crossing          183142 non-null float64
finishing         183142 non-null float64
heading_accuracy  183142 non-null float64
short_passing     183142 non-null float64
volleys           181265 non-null float64
dribbling         183142 non-null float64
curve             181265 non-null float64
free_kick_accuracy 183142 non-null float64
long_passing      183142 non-null float64
ball_control      183142 non-null float64
acceleration      183142 non-null float64
sprint_speed      183142 non-null float64
agility           181265 non-null float64
reactions         183142 non-null float64
balance           181265 non-null float64
shot_power        183142 non-null float64
jumping           181265 non-null float64
stamina           183142 non-null float64
strength          183142 non-null float64
long_shots        183142 non-null float64
aggression        183142 non-null float64
interceptions     183142 non-null float64
positioning       183142 non-null float64
vision            181265 non-null float64
penalties         183142 non-null float64
marking           183142 non-null float64
standing_tackle   183142 non-null float64
sliding_tackle    181265 non-null float64
gk_diving         183142 non-null float64
gk_handling       183142 non-null float64
gk_kicking        183142 non-null float64
gk_positioning    183142 non-null float64
gk_reflexes       183142 non-null float64
dtypes: float64(35), int64(3), object(4)
memory usage: 56.1+ MB
```

```
In [7]: df = df.dropna()
```

In [8]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 180354 entries, 0 to 183977
Data columns (total 42 columns):
id                180354 non-null int64
player_fifa_api_id 180354 non-null int64
player_api_id     180354 non-null int64
date              180354 non-null object
overall_rating     180354 non-null float64
potential          180354 non-null float64
preferred_foot     180354 non-null object
attacking_work_rate 180354 non-null object
defensive_work_rate 180354 non-null object
crossing           180354 non-null float64
finishing          180354 non-null float64
heading_accuracy   180354 non-null float64
short_passing      180354 non-null float64
volleys            180354 non-null float64
dribbling          180354 non-null float64
curve              180354 non-null float64
free_kick_accuracy 180354 non-null float64
long_passing       180354 non-null float64
ball_control       180354 non-null float64
acceleration       180354 non-null float64
sprint_speed       180354 non-null float64
agility            180354 non-null float64
reactions          180354 non-null float64
balance            180354 non-null float64
shot_power         180354 non-null float64
jumping            180354 non-null float64
stamina            180354 non-null float64
strength           180354 non-null float64
long_shots         180354 non-null float64
aggression         180354 non-null float64
interceptions      180354 non-null float64
positioning        180354 non-null float64
vision             180354 non-null float64
penalties          180354 non-null float64
marking            180354 non-null float64
standing_tackle    180354 non-null float64
sliding_tackle     180354 non-null float64
gk_diving          180354 non-null float64
gk_handling        180354 non-null float64
gk_kicking         180354 non-null float64
gk_positioning     180354 non-null float64
gk_reflexes        180354 non-null float64
dtypes: float64(35), int64(3), object(4)
memory usage: 56.4+ MB
```

In [9]: *# Feature selection - Since we have about 40 different features, we can run some feature selection algorithms to reduce the size of our featureset.*

```
from sklearn.preprocessing import scale
from sklearn.feature_selection import RFE
```

```
In [10]: df_new = df.copy()
```

```
In [11]: # Most of our data is numeric, with few exceptions having floating data types.  
In the subsequent prediction analysis  
# we'll only concern ourself with the integer numerics, but there is obviously  
potential gains to be made by  
# incorporating the qualitative data (i.e. player position).  
df_new = df_new.select_dtypes(["int64", "float64"])
```

```
In [12]: df_new.shape
```

```
Out[12]: (180354, 38)
```

In [13]: df_new.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 180354 entries, 0 to 183977
Data columns (total 38 columns):
id                180354 non-null int64
player_fifa_api_id 180354 non-null int64
player_api_id     180354 non-null int64
overall_rating     180354 non-null float64
potential         180354 non-null float64
crossing          180354 non-null float64
finishing         180354 non-null float64
heading_accuracy  180354 non-null float64
short_passing     180354 non-null float64
volleys          180354 non-null float64
dribbling         180354 non-null float64
curve            180354 non-null float64
free_kick_accuracy 180354 non-null float64
long_passing      180354 non-null float64
ball_control      180354 non-null float64
acceleration      180354 non-null float64
sprint_speed      180354 non-null float64
agility           180354 non-null float64
reactions         180354 non-null float64
balance           180354 non-null float64
shot_power        180354 non-null float64
jumping           180354 non-null float64
stamina           180354 non-null float64
strength          180354 non-null float64
long_shots        180354 non-null float64
aggression        180354 non-null float64
interceptions     180354 non-null float64
positioning       180354 non-null float64
vision            180354 non-null float64
penalties         180354 non-null float64
marking           180354 non-null float64
standing_tackle   180354 non-null float64
sliding_tackle    180354 non-null float64
gk_diving         180354 non-null float64
gk_handling       180354 non-null float64
gk_kicking        180354 non-null float64
gk_positioning    180354 non-null float64
gk_reflexes       180354 non-null float64
dtypes: float64(35), int64(3)
memory usage: 53.7 MB
```

In [14]: `X = df_new.drop('overall_rating',axis=1).values`
`y = df_new['overall_rating'].values.ravel()`
`from sklearn.preprocessing import scale`
`X = scale(X)`

In [15]: `df_1 = df_new.drop('overall_rating',axis=1)`

In [16]: X.shape

Out[16]: (180354, 37)

In [17]: y.shape

Out[17]: (180354,)

In [18]: *# Feature Selection using RFE Scikit Library*

```
lm = LinearRegression()

rfe = RFE(lm, n_features_to_select = 10)
rfe_fit = rfe.fit(X, y)
features = []
for feat in df_1.columns[rfe_fit.support_]:
    print(feat)
    features.append(feat)
```

```
player_api_id
potential
heading_accuracy
short_passing
ball_control
reactions
strength
gk_diving
gk_kicking
gk_positioning
```

In [19]: features

Out[19]: ['player_api_id',
'potential',
'heading_accuracy',
'short_passing',
'ball_control',
'reactions',
'strength',
'gk_diving',
'gk_kicking',
'gk_positioning']

```
In [20]: # Using Statsmodel to illustrate the summary results
lm = LinearRegression()

rfe = RFE(lm, n_features_to_select = 15)
rfe_fit = rfe.fit(X, y)
features = []
for feat in df_1.columns[rfe_fit.support_]:
    print(feat)
    features.append(feat)

player_fifa_api_id
player_api_id
potential
heading_accuracy
short_passing
ball_control
acceleration
reactions
strength
marking
gk_diving
gk_handling
gk_kicking
gk_positioning
gk_reflexes
```

```
In [21]: df_optm = df_new[features]
```

```
In [22]: df_optm.shape
```

```
Out[22]: (180354, 15)
```



```
In [23]: # Using Statsmodels for analysing the impact of attribute potential on the player rating
import statsmodels.formula.api as sm
model1 = sm.OLS(df_new['overall_rating'], df_new['potential'])
result1 = model1.fit()
print(result1.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          overall_rating    R-squared:                0.99
6
Model:                  OLS              Adj. R-squared:            0.99
6
Method:                 Least Squares    F-statistic:                 4.062e+0
7
Date:                  Wed, 06 Mar 2019   Prob (F-statistic):          0.0
0
Time:                  12:26:23          Log-Likelihood:              -5.3063e+0
5
No. Observations:      180354           AIC:                        1.061e+0
6
Df Residuals:          180353           BIC:                        1.061e+0
6
Df Model:              1
Covariance Type:       nonrobust
=====
```

```
=====
              coef      std err          t      P>|t|      [0.025      0.97
5]
-----
potential      0.9331      0.000    6373.716      0.000      0.933      0.93
3
=====
```

```
=====
Omnibus:              30950.684    Durbin-Watson:              0.37
5
Prob(Omnibus):         0.000    Jarque-Bera (JB):            54474.35
7
Skew:                  -1.111    Prob(JB):                     0.0
0
Kurtosis:              4.520    Cond. No.                     1.0
0
=====
```

```
=====
=
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```



```
In [24]: # Using Statsmodels for analysing the impact of all attribute on the player rating  
X_new = df_new[features].values  
model = sm.OLS(df_new['overall_rating'],df_new[features])  
result = model.fit()  
print(result.summary())
```

OLS Regression Results

```

=====
=
Dep. Variable:          overall_rating    R-squared:                0.99
9
Model:                  OLS              Adj. R-squared:          0.99
9
Method:                 Least Squares    F-statistic:             8.035e+0
6
Date:                   Wed, 06 Mar 2019  Prob (F-statistic):       0.0
0
Time:                   12:27:20         Log-Likelihood:          -4.3282e+0
5
No. Observations:       180354          AIC:                     8.657e+0
5
Df Residuals:           180339          BIC:                     8.658e+0
5
Df Model:               15
Covariance Type:        nonrobust
=====
=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
player_fifa_api_id -8.998e-06   1.44e-07   -62.432    0.000   -9.28e-06
-8.72e-06
player_api_id      -6.663e-06   6.01e-08  -110.802    0.000   -6.78e-06
-6.55e-06
potential           0.4525      0.001    327.126    0.000      0.450
0.455
heading_accuracy    0.0643      0.001     93.120    0.000      0.063
0.066
short_passing       0.0646      0.001     59.132    0.000      0.062
0.067
ball_control        0.1317      0.001    104.742    0.000      0.129
0.134
acceleration        0.0299      0.001     43.957    0.000      0.029
0.031
reactions           0.1811      0.001    175.817    0.000      0.179
0.183
strength            0.0600      0.001     91.087    0.000      0.059
0.061
marking             0.0187      0.000     49.000    0.000      0.018
0.019
gk_diving           0.1627      0.001    117.671    0.000      0.160
0.165
gk_handling         0.0254      0.002     13.865    0.000      0.022
0.029
gk_kicking          -0.0479     0.001    -83.797    0.000     -0.049
-0.047
gk_positioning      0.0437      0.002     23.984    0.000      0.040
0.047
gk_reflexes         0.0179      0.002      9.987    0.000      0.014
0.021
=====
=

```

Omnibus:	13814.579	Durbin-Watson:	0.42
5			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	34846.32
8			
Skew:	-0.457	Prob(JB):	0.0
0			
Kurtosis:	4.950	Cond. No.	8.42e+0
4			

=====

=

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 8.42e+04. This might indicate that there are strong multicollinearity or other numerical problems.




In [25]:

```
# Explanation of the OLS Regression Results :
# Adjusted R-squared indicates that 99.9% of player ratings can be explained by
# our predictor variable.
# The regression coefficient (coef) represents the change in the dependent variable
# resulting from a one unit change in
# the predictor variable, all other variables being held constant.
# In our model, a one unit increase in potential increases the rating by 0.4525.
# The standard error measures the accuracy of potential's coefficient by estimating
# the variation of the coefficient
# if the same test were run on a different sample of our population.
# Our standard error, 0.001, is low and therefore appears accurate.
# The p-value means the probability of an 0.4525 increasing in player rating due
# to a one unit increase in potential
# is 0%, assuming there is no relationship between the two variables.
# A low p-value indicates that the results are statistically significant, that is
# in general the p-value is less than 0.05.
# The confidence interval is a range within which our coefficient is likely to
# fall. We can be 95% confident that
# potential's coefficient will be within our confidence interval, [0.450, 0.455].
# Warnings in the Summary provided by statsmodels OLS model
# Multicollinearity: A careful observer would've noticed the warnings produced
# by our model regarding multicollinearity.
# We have two or more variables telling roughly the same story, overstating the
# value of each of the predictors.

# REGRESSION PLOTS

import matplotlib.pyplot as plt
import statsmodels.api as sm
%matplotlib inline
fig = plt.figure(figsize=(20,12))
fig = sm.graphics.plot_partregress_grid(result, fig=fig)
```



```
In [26]: import statsmodels.formula.api as smf
# only include TV and Radio in the model
lm = smf.ols(formula='overall_rating ~ player_fifa_api_id + player_api_id + potential + heading_accuracy + short_passing + ball_control + acceleration + reaction_s + strength + marking + gk_diving + gk_handling + gk_kicking + gk_positioning + gk_reflexes ', data=df_new).fit()
print('Confidence of the statsmodel for the input data : ',lm.rsquared)
```


Confidence of the statsmodel for the input data : 0.8561047414083593

```
In [27]: df_new.head()
```

Out[27]:

	id	player_fifa_api_id	player_api_id	overall_rating	potential	crossing	finishing	heading_accu
0	1	218353	505942	67.0	71.0	49.0	44.0	
1	2	218353	505942	67.0	71.0	49.0	44.0	
2	3	218353	505942	62.0	66.0	49.0	44.0	
3	4	218353	505942	61.0	65.0	48.0	43.0	
4	5	218353	505942	61.0	65.0	48.0	43.0	


5 rows × 38 columns



```
In [28]: df_optm.head()
```

Out[28]:

	player_fifa_api_id	player_api_id	potential	heading_accuracy	short_passing	ball_control	acc
0	218353	505942	71.0	71.0	61.0	49.0	
1	218353	505942	71.0	71.0	61.0	49.0	
2	218353	505942	66.0	71.0	61.0	49.0	
3	218353	505942	65.0	70.0	60.0	48.0	
4	218353	505942	65.0	70.0	60.0	48.0	




```
In [29]: # Data Exploration using visualization
df_corr = df_new.corr()
```

```
In [30]: print('The features contain high correlation .We need to remove them first before
re applying Regression Techniques.')
import seaborn as sns
sns.set_style('whitegrid')
plt.figure(figsize=(20,8))
sns.heatmap(df_corr,annot=True)
```

The features contain high correlation .We need to remove them first before applying Regression Techniques.

```
Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0x107cef30>
```



```
In [31]: #create correlation matrix with absolute values
df_corr = df_new.corr().abs()
#select upper triangle of matrix
up_tri = df_corr.where(np.triu(np.ones(df_corr.shape[1]),k=1).astype(np.bool))

#find all the features which have a correlation > 0.75 with other features.
corr_features = [ column for column in up_tri.columns if any(up_tri[column]>
0.75)]

#Print Correlated features
print(corr_features)
```

```
['potential', 'short_passing', 'volleys', 'dribbling', 'curve', 'free_kick_ac
curacy', 'long_passing', 'ball_control', 'sprint_speed', 'agility', 'reaction
s', 'shot_power', 'long_shots', 'positioning', 'vision', 'penalties', 'markin
g', 'standing_tackle', 'sliding_tackle', 'gk_handling', 'gk_positioning', 'gk
_reflexes']
```

```
In [32]: #Drop Correlated Features  
df_no_corr = df_new.drop(corr_features,axis=1)  
df_no_corr.head()
```

```
Out[32]:
```

	id	player_fifa_api_id	player_api_id	overall_rating	crossing	finishing	heading_accuracy	acc
0	1	218353	505942	67.0	49.0	44.0	71.0	
1	2	218353	505942	67.0	49.0	44.0	71.0	
2	3	218353	505942	62.0	49.0	44.0	71.0	
3	4	218353	505942	61.0	48.0	43.0	70.0	
4	5	218353	505942	61.0	48.0	43.0	70.0	


```
In [33]: # This shows that the feature selection API - sklearn.feature_selection.RFE has  
resulted in the same feature  
# selection for top 15 features selected.  
# Quant Features against Rating  
  
len(df_no_corr.columns)
```

```
Out[33]: 16
```



```
In [34]: import matplotlib.pyplot as plt
fig = plt.figure(figsize = (15,60))
val = df_optm.shape[1]
for idx in range(val):
    feature = df_optm.columns[idx]
    ax = fig.add_subplot(13,3,idx+1)
    Xtmp = df_optm[feature]
    ax.scatter(Xtmp, y)
    ax.set_xlabel(feature)

plt.tight_layout()
plt.show()
```



```
In [35]: # Split the input data into training and test data
from sklearn.model_selection import train_test_split
#splitting 66.66% for train data and 33.33% for test data.
X_train,X_test,Y_train,Y_test = train_test_split(X,y,test_size=0.33,random_state=0)
print("X_train Shape : ",X_train.shape)
print("X_test Shape : ",X_test.shape)
print("Y_train Shape : ",Y_train.shape)
print("Y_test.shape : ",Y_test.shape)
```

```
X_train Shape : (120837, 37)
X_test Shape : (59517, 37)
Y_train Shape : (120837,)
Y_test.shape : (59517,)
```

```
In [36]: # Applying Linear Regression Model
lm = LinearRegression()
lm.fit(X_train, Y_train)# train the model
```

```
Out[36]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```


```
In [37]: # Perform Prediction using Linear Regression Model
Y_pred = lm.predict(X_test) # predict the prices based on the test data
Y_pred
```

```
Out[37]: array([73.74672865, 68.65401302, 67.22132467, ..., 71.6247244 ,
64.76934151, 62.18654694])
```

```
In [38]: print("The variance score of the LinearRegression model is : ",lm.score(X_test,Y_test))
print('Since variance score is near about 1 it seems to be a perfect prediction')
```


```
The variance score of the LinearRegression model is : 0.8593275836597539
Since variance score is near about 1 it seems to be a perfect prediction
```

```
In [39]: import matplotlib.pyplot as plt
plt.figure(figsize=(6, 4))
plt.hist(Y_pred)
plt.xlabel('Predicted Rating of the Player')
plt.ylabel('count')
plt.tight_layout()
```



```
In [40]: import seaborn as sns
sns.set_style('whitegrid')
plt.figure(figsize=(10, 8))
plt.scatter(Y_test, Y_pred)
plt.plot([0, 100], [0, 100], '--k')
plt.axis('tight')
plt.xlabel('Actual Rating')
plt.ylabel('Predicted Rating')
plt.tight_layout()
plt.title("Actual vs Predicted Rating for LinearRegression Model")
```

Out[40]: Text(0.5, 1.0, 'Actual vs Predicted Rating for LinearRegression Model')



```
In [41]: # Evaluate Linear Regression Accuracy using Root Mean Square Error
from sklearn.metrics import mean_squared_error
print("Error Rate of the Regression Model : ", sqrt(mean_squared_error(Y_pred, Y_test)))
```

Error Rate of the Regression Model : 2.623473911844569

```
In [42]: # Applying Decision Tree Regressor Model to the input data
regressor = DecisionTreeRegressor(max_depth=20)
regressor.fit(X_train, Y_train)
```

```
Out[42]: DecisionTreeRegressor(criterion='mse', max_depth=20, max_features=None,
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
presort=False, random_state=None, splitter='best')
```

```
In [43]: # Perform Prediction using Decision Tree Regressor

Y_pred = regressor.predict(X_test)
Y_pred
```


```
Out[43]: array([[76.          , 72.          , 67.          , ..., 71.          ,
62.          , 61.78571429])
```

```
In [44]: print("The variance score of the DecisionTreeRegressor model is :",regressor
.score(X_test,Y_test))
print('Since variance score is near about 1 it seems to be a perfect predictio
n')
```

The variance score of the DecisionTreeRegressor model is : 0.9547509358789791


Since variance score is near about 1 it seems to be a perfect prediction

```
In [45]: import matplotlib.pyplot as plt
plt.figure(figsize=(6, 4))
plt.hist(Y_pred)
plt.xlabel('Predicted Rating of the Player')
plt.ylabel('count')
plt.tight_layout()
```



```
In [46]: plt.figure(figsize=(10, 8))
plt.scatter(Y_test, Y_pred)
plt.plot([0, 100], [0, 100], '--k')
plt.axis('tight')
plt.xlabel('Actual Rating')
plt.ylabel('Predicted Rating')
plt.tight_layout()
plt.title("Actual vs Predicted Rating for DecisionTreeRegressor Model")
```

Out[46]: Text(0.5, 1.0, 'Actual vs Predicted Rating for DecisionTreeRegressor Model')



```
In [47]: #The mean of the expected target value in test set
Y_test.mean()
```

Out[47]: 68.66518809751835

```
In [48]: #The mean of the predicted target value in test set ?
Y_pred.mean()
```

Out[48]: 68.65974748152661

```
In [49]: # Evaluate Linear Regression Accuracy using Root Mean Square Error For DecisionTreeRegressor model
print("Error Rate of the DecisionTreeRegressor Model : ",sqrt(mean_squared_error(Y_pred,Y_test)))
print('The DecisionTreeRegressor Model performs better than the LinearRegression Model as evident from the error rate')
```

Error Rate of the DecisionTreeRegressor Model : 1.487911418908021
 The DecisionTreeRegressor Model performs better than the LinearRegression Model as evident from the error rate


```
In [50]: # Obtaining predictions by cross-validation for the Regression Models
```

```
df_optm = df_new.copy()
df_optm['rating'] = y
df_optm.head()
```

Out[50]:

	id	player_fifa_api_id	player_api_id	overall_rating	potential	crossing	finishing	heading_accuracy
0	1	218353	505942	67.0	71.0	49.0	44.0	
1	2	218353	505942	67.0	71.0	49.0	44.0	
2	3	218353	505942	62.0	66.0	49.0	44.0	
3	4	218353	505942	61.0	65.0	48.0	43.0	
4	5	218353	505942	61.0	65.0	48.0	43.0	

5 rows × 39 columns



```
In [51]: from sklearn.model_selection import cross_val_predict
X = df_optm.drop('rating',axis=1)
Y = df_optm['rating']
predicted = cross_val_predict(regressor, X, Y, cv=10)
```

```
In [52]: from sklearn.metrics import accuracy_score
print( "Accuracy Score of the DecisionTreeRegressor Model is : " ,accuracy_score(Y.astype(int), predicted) )
```

Accuracy Score of the DecisionTreeRegressor Model is : 0.9997615800037704

```
In [62]: # Calculate Error using K-Fold Cross validation
from sklearn.cross_validation import KFold
kfold = KFold(len(df_optm),n_folds=10,shuffle=True,random_state=0)
```



```
In [65]: from sklearn.metrics import mean_absolute_error
lm = LinearRegression()
mean_abs_error = []
accuracy_score = []
for train,test in kfold:
    x = X.iloc[train]
    y = Y.iloc[train]
    lm.fit(x,y)
    Y_test = Y.iloc[test]
    Y_pred = lm.predict(X.iloc[test])
    mean_abs_error.append(mean_absolute_error(Y_test,Y_pred))
```

```
In [70]: print('10 Fold Cross Validation Error : {} accuracy score : {} for LinearRegression Model '.format(np.mean(mean_abs_error),1 - np.mean(mean_abs_error)))
```

10 Fold Cross Validation Error : 3.659946305405014e-12 accuracy score : 0.9999999999634 for LinearRegression Model

```
In [71]: from sklearn.metrics import mean_absolute_error
#DR = LinearRegression()
mean_abs_error = []
accuracy_score = []
for train,test in kfold:
    x = X.iloc[train]
    y = Y.iloc[train]
    regressor.fit(x,y)
    Y_test = Y.iloc[test]
    Y_pred = regressor.predict(X.iloc[test])
    mean_abs_error.append(mean_absolute_error(Y_test,Y_pred))
```

```
In [72]: print('10 Fold Cross Validation Error : {} accuracy score : {} for DecisionTreeRegressor Model '.format(np.mean(mean_abs_error),1 - np.mean(mean_abs_error)))
```

10 Fold Cross Validation Error : 4.435819240365955e-05 accuracy score : 0.9999556418075963 for DecisionTreeRegressor Model

```
In [ ]: # We have use the below models to predict the player ratings:
# 1) Statsmodels.api.OLS
# 2) LinearRegression
# 3) DecisionTreeRegressor

# Sampling Mechanisms used:
# 1) Test Train Split
# 2) 10 Fold Cross Validation

# Model Estimation mechanisms used:
# 1) Root Mean Squared Error
# 2) 10 Fold Cross Validation error.
```